

# Introduction to Machine Learning

## Neural Networks

Varun Chandola

Computer Science & Engineering  
State University of New York at Buffalo  
Buffalo, NY, USA  
chandola@buffalo.edu



University at Buffalo  
Department of Computer Science  
and Engineering  
School of Engineering and Applied Sciences



## Extending Perceptrons

## Multi Layered Perceptrons

- Generalizing to Multiple Labels

- Properties of Sigmoid Function

- Motivation for Using Non-linear Surfaces

## Feed Forward Neural Networks

## Backpropagation

- Derivation of the Backpropagation Rules


## Final Algorithm

## Wrapping up Neural Networks

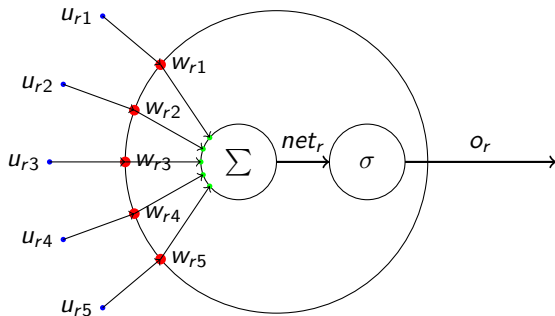
## Bias Variance Tradeoff

# Extending Perceptrons

- ▶ Questions?
  - ▶ Why not work with thresholded perceptron?
    - ▶ Not differentiable
  - ▶ How to learn non-linear surfaces?
  - ▶ How to generalize to multiple outputs, numeric output?



# Anatomy of a Sigmoid Unit ( $r$ )

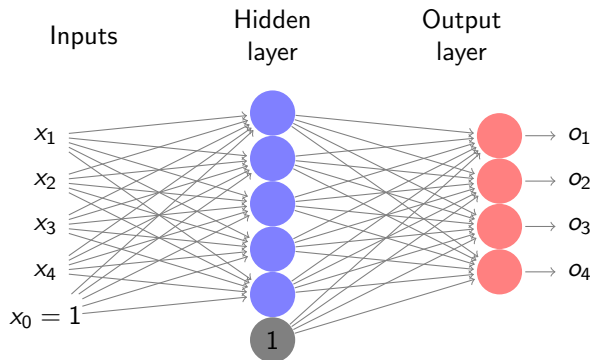


# Generalizing to Multiple Labels

- ▶ Distinguishing between multiple categories
- ▶ *Solution:* Add another layer - **Multi Layer Neural Networks**

# Generalizing to Multiple Labels

- ▶ Distinguishing between multiple categories
- ▶ *Solution:* Add another layer - **Multi Layer Neural Networks**



# What Threshold Unit to Use?

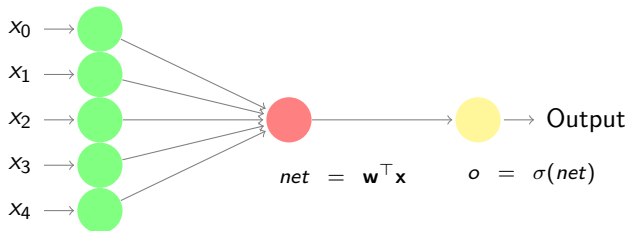
- ▶ Linear Unit
- ▶ Perceptron Unit

# What Threshold Unit to Use?

- ▶ Linear Unit
- ▶ Perceptron Unit
- ▶ Sigmoid Unit
  - ▶ Smooth, differentiable threshold function

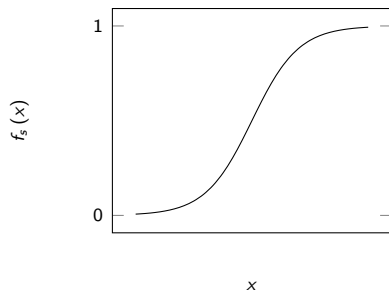
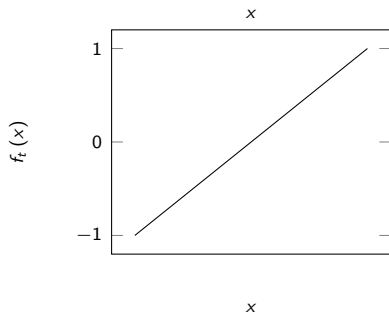
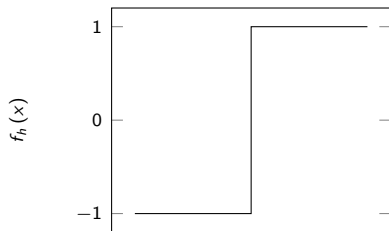
$$\sigma(net) = \frac{1}{1 + e^{-net}}$$

- ▶ Non-linear output

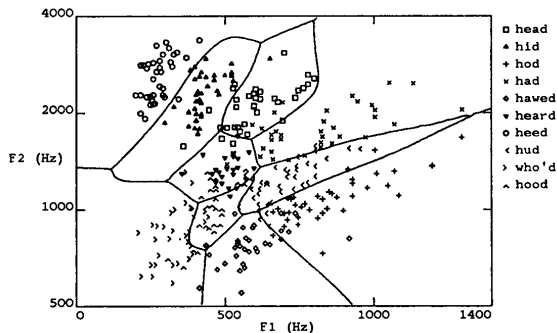





# Properties of Sigmoid Function



# Motivation for Using Non-linear Surfaces



# Feed Forward Neural Networks - Architecture



# Feed Forward Neural Networks

- ▶  $D$  input nodes (excluding bias)
  - ▶  $M$  hidden nodes (excluding bias)
  - ▶  $K$  output nodes
- ▶ At hidden nodes:  $\mathbf{w}_j, 1 \leq j \leq M, \mathbf{w}_j \in \mathbb{R}^{D+1}$
  - ▶ At output nodes:  $\mathbf{w}_l, 1 \leq l \leq K, \mathbf{w}_l \in \mathbb{R}^{M+1}$

# Learning Weights of the Multi-layer Network

- ▶ Assume that the network structure is predetermined (number of hidden nodes and interconnections)
- ▶ Objective function for  $N$  training examples:

$$J = \sum_{i=1}^N J_i = \frac{1}{2} \sum_{i=1}^N \sum_{l=1}^K (y_{il} - o_{il})^2$$

- ▶  $y_{il}$  - Target value associated with  $l^{th}$  class for input ( $\mathbf{x}_i$ )
- ▶  $y_{il} = 1$  when  $k$  is true class for  $\mathbf{x}_i$ , and 0 otherwise
- ▶  $o_{il}$  - Predicted output value at  $l^{th}$  output node for  $\mathbf{x}_i$

What are we learning?

Weight vectors for all output and hidden nodes that minimize  $J$

# The Backpropagation Algorithm

1. Initialize all weights to *small values*
2. For each training example,  $\langle \mathbf{x}, \mathbf{y} \rangle$ :
  - 2.1 **Propagate input forward** through the network
  - 2.2 **Propagate errors backward** through the network

# Backpropagation Algorithm - Continued

## Gradient Descent

- ▶ Move in the opposite direction of the **gradient** of the objective function
- ▶  $-\eta \nabla J$

$$\nabla J = \sum_{i=1}^N \nabla J_i$$

- ▶ What is the gradient computed with respect to?

- ▶ Weights -  $m$  at hidden nodes and  $k$  at output nodes

- ▶  $\mathbf{w}_j$  ( $j = 1 \dots m$ )

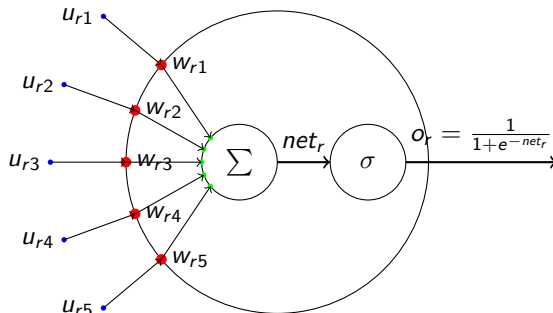
- ▶  $\mathbf{w}_l$  ( $l = 1 \dots k$ )

- ▶  $\mathbf{w}_j \leftarrow \mathbf{w}_j - \eta \frac{\partial J}{\partial \mathbf{w}_j} = \mathbf{w}_j - \eta \sum_{i=1}^N \frac{\partial J_i}{\partial \mathbf{w}_j}$

- ▶  $\mathbf{w}_l \leftarrow \mathbf{w}_l - \eta \frac{\partial J}{\partial \mathbf{w}_l} = \mathbf{w}_l - \eta \sum_{i=1}^N \frac{\partial J_i}{\partial \mathbf{w}_l}$

$$\nabla J_i = \begin{bmatrix} \frac{\partial J_i}{\partial \mathbf{w}_1} \\ \frac{\partial J_i}{\partial \mathbf{w}_2} \\ \vdots \\ \frac{\partial J_i}{\partial \mathbf{w}_{m+k}} \end{bmatrix}$$

# Anatomy of a Sigmoid Unit ( $r$ )



$$\frac{\partial J_i}{\partial \mathbf{w}_r} = \begin{bmatrix} \frac{\partial J_i}{\partial w_{r1}} \\ \frac{\partial J_i}{\partial w_{r2}} \\ \vdots \end{bmatrix}$$

- ▶ Need to compute  $\frac{\partial J_i}{\partial w_{rq}}$
- ▶ Update rule for the  $q^{th}$  entry in the  $r^{th}$  weight vector:

$$w_{rq} \leftarrow w_{rq} - \eta \frac{\partial J}{\partial w_{rq}} = w_{rq} - \eta \sum_{i=1}^N \frac{\partial J_i}{\partial w_{rq}}$$



# Derivation of the Backpropagation Rules

Assume that we only one training example, i.e.,  $i = 1$ ,  $J = J_i$ . Dropping the subscript  $i$  from here onwards.

- ▶ Consider any weight  $w_{rq}$
- ▶ Let  $u_{rq}$  be the  $q^{th}$  element of the input vector coming in to the  $r^{th}$  unit.

## Observation 1

Weight  $w_{rq}$  is connected to  $J$  through  $net_r = \sum_i w_{rq} u_{rq}$ .

$$\frac{\partial J}{\partial w_{rq}} = \frac{\partial J}{\partial net_r} \frac{\partial net_r}{\partial w_{rq}} = \frac{\partial J}{\partial net_r} u_{rq}$$

# Analyzing Output Nodes

## Observation 2

$net_l$  for an **output node** is connected to  $J$  only through the output value of the node (or  $o_l$ )

$$\frac{\partial J}{\partial net_l} = \frac{\partial J}{\partial o_l} \frac{\partial o_l}{\partial net_l}$$

# Analyzing Output Nodes

## Observation 2

$net_I$  for an **output node** is connected to  $J$  only through the output value of the node (or  $o_I$ )

$$\frac{\partial J}{\partial net_I} = \frac{\partial J}{\partial o_I} \frac{\partial o_I}{\partial net_I}$$

## Update Rule for Output Units

$$w_{Ij} \leftarrow w_{Ij} + \eta \delta_I u_{Ij}$$

where  $\delta_I = (y_I - o_I)o_I(1 - o_I)$ .

► *Question:* What is  $u_{Ij}$  for the  $I^{th}$  output node?

# Analyzing Hidden Nodes

## Observation 3

$net_j$  for a **hidden node** is connected to  $J$  through all output nodes

$$\frac{\partial J}{\partial net_j} = \sum_{l=1}^K \frac{\partial J}{\partial net_l} \frac{\partial net_l}{\partial net_j}$$

# Analyzing Hidden Nodes

## Observation 3

$net_j$  for a **hidden node** is connected to  $J$  through all output nodes

$$\frac{\partial J}{\partial net_j} = \sum_{l=1}^K \frac{\partial J}{\partial net_l} \frac{\partial net_l}{\partial net_j}$$

## Update Rule for Hidden Units

$$w_{jp} \leftarrow w_{jp} + \eta \delta_j u_{jp}$$

# Analyzing Hidden Nodes

## Observation 3

$net_j$  for a **hidden node** is connected to  $J$  through all output nodes

$$\frac{\partial J}{\partial net_j} = \sum_{l=1}^K \frac{\partial J}{\partial net_l} \frac{\partial net_l}{\partial net_j}$$

## Update Rule for Hidden Units

$$w_{jp} \leftarrow w_{jp} + \eta \delta_j u_{jp}$$

$$\delta_j = o_j(1 - o_j) \sum_{l=1}^K \delta_l w_{lj}$$

$$\delta_l = (y_l - o_l) o_l(1 - o_l)$$

► *Question:* What is  $u_{jp}$  for the  $j^{th}$  hidden node?

# Final Algorithm

- ▶ While not converged:
  - ▶ *Move forward* to compute outputs at hidden and output nodes
  - ▶ *Move backward* to propagate errors back
    - ▶ Compute  $\delta$  errors at output nodes ( $\delta_l$ )
    - ▶ Compute  $\delta$  errors at hidden nodes ( $\delta_j$ )
  - ▶ Update all weights according to weight update equations

# Conclusions about Neural Networks

- ▶ Error function contains many local minima
- ▶ No guarantee of convergence
  - ▶ Not a “big” issue in practical deployments
- ▶ Improving backpropagation



# Conclusions about Neural Networks

- ▶ Error function contains many local minima
- ▶ No guarantee of convergence
  - ▶ Not a “big” issue in practical deployments
- ▶ Improving backpropagation
  - ▶ Adding momentum
  - ▶ Using stochastic gradient descent
  - ▶ Train multiple times using different initializations

# Bias Variance Tradeoff

- ▶ Neural networks are *universal function approximators*
  - ▶ By making the model more complex (increasing number of hidden layers or  $m$ ) one can lower the error
- ▶ Is the model with least training error the best model?

# Bias Variance Tradeoff

- ▶ Neural networks are *universal function approximators*
  - ▶ By making the model more complex (increasing number of hidden layers or  $m$ ) one can lower the error
- ▶ Is the model with least training error the best model?
  - ▶ The simple answer is **no!**
  - ▶ Risk of overfitting (chasing the data)
  - ▶ Overfitting  $\Leftarrow$  **High generalization error**

## High Variance - Low Bias

- ▶ “Chases the data”
- ▶ Very low training error
- ▶ Poor performance on unseen data

## Low Variance - High Bias

- ▶ Less sensitive to training data
- ▶ Higher training error
- ▶ Better performance on unseen data

# Getting the Right Balance

- ▶ General rule of thumb – If two models are giving similar training error, choose the **simpler** model
- ▶ What is simple for a neural network?
- ▶ Low weights in the weight matrices?
  - ▶ Why?

# Introducing Bias in Neural Network Training

- ▶ Penalize solutions in which the weights are high
- ▶ Can be done by introducing a penalty term in the objective function
  - ▶ **Regularization**

## Regularization for Backpropagation

$$\tilde{J} = J + \frac{\lambda}{2n} \left( \sum_{j=1}^M \sum_{i=1}^{D+1} (w_{ji}^{(1)})^2 + \sum_{l=1}^K \sum_{j=1}^{M+1} (w_{lj}^{(2)})^2 \right)$$

# Other Extensions?

- ▶ Use a different loss function (why)?
  - ▶ Quadratic (Squared), Cross-entropy, Exponential, KL Divergence, etc.
- ▶ Use a different activation function (why)?
  - ▶ Sigmoid

$$f(z) = \frac{1}{1 + \exp(-z)}$$

- ▶ Tanh

$$f(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

- ▶ Rectified Linear Unit (ReLU)

$$f(z) = \max(0, z)$$

# References